

Introduction to EDA

Pramod Duvvuri

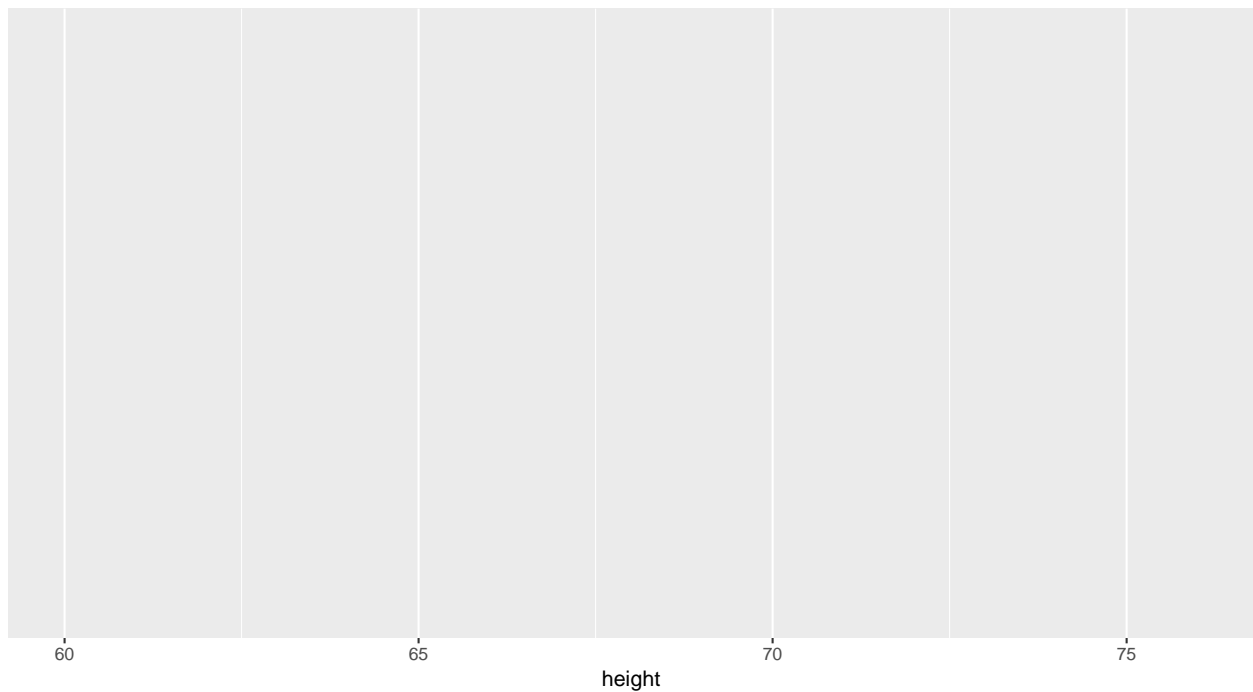
3/11/2019

Univariate Data

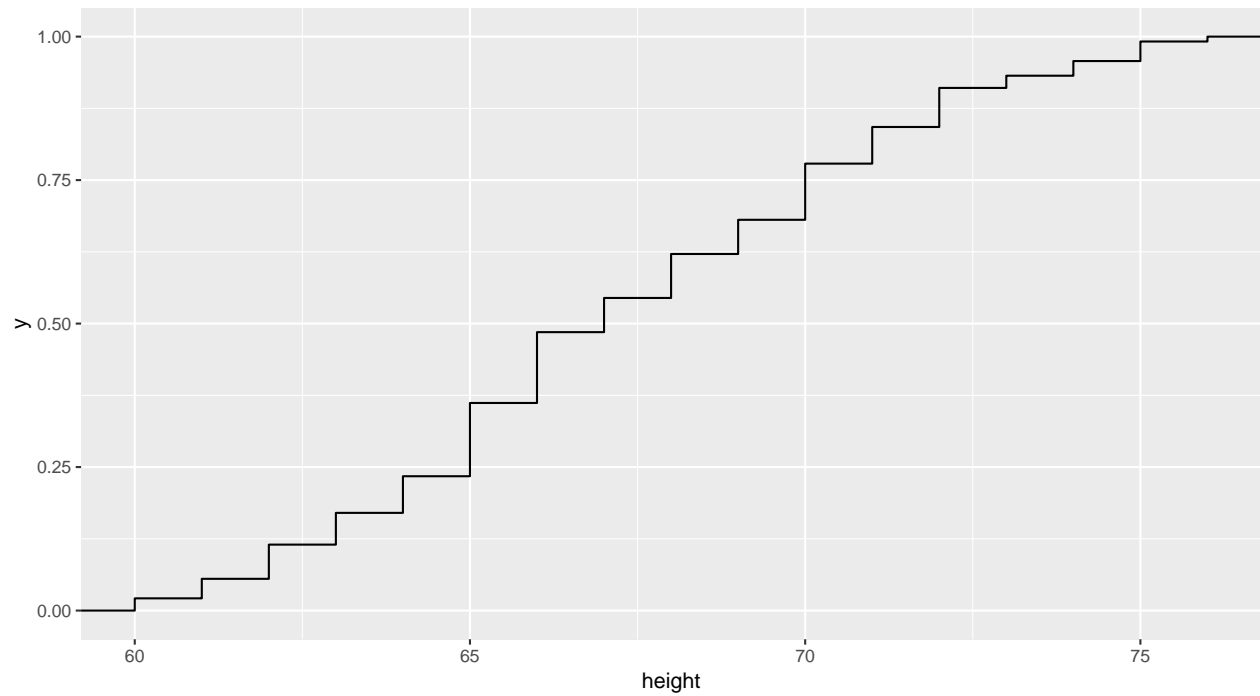
Learning Grammar of Graphics (ggplot2)

```
# Load installed packages
library('lattice')
library('ggplot2')

ggplot(singer, aes(x = height))
```

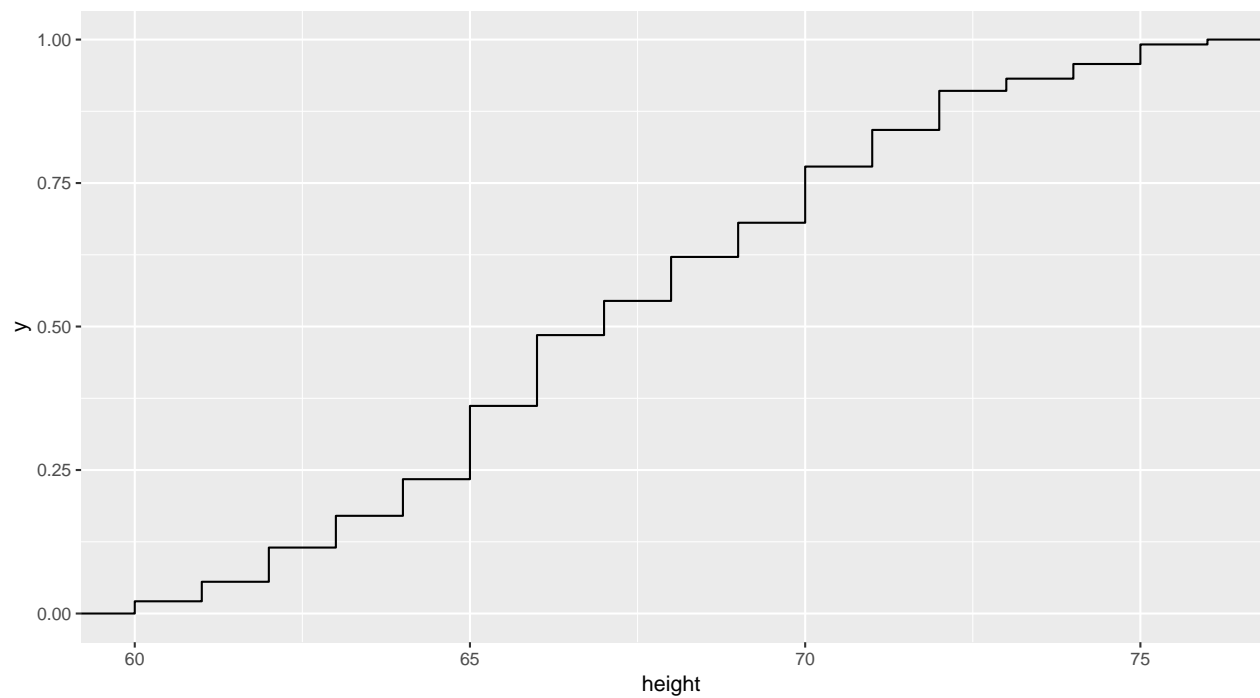


```
# ECDF in ggplot2
ggplot(singer, aes(x = height)) + stat_ecdf()
```

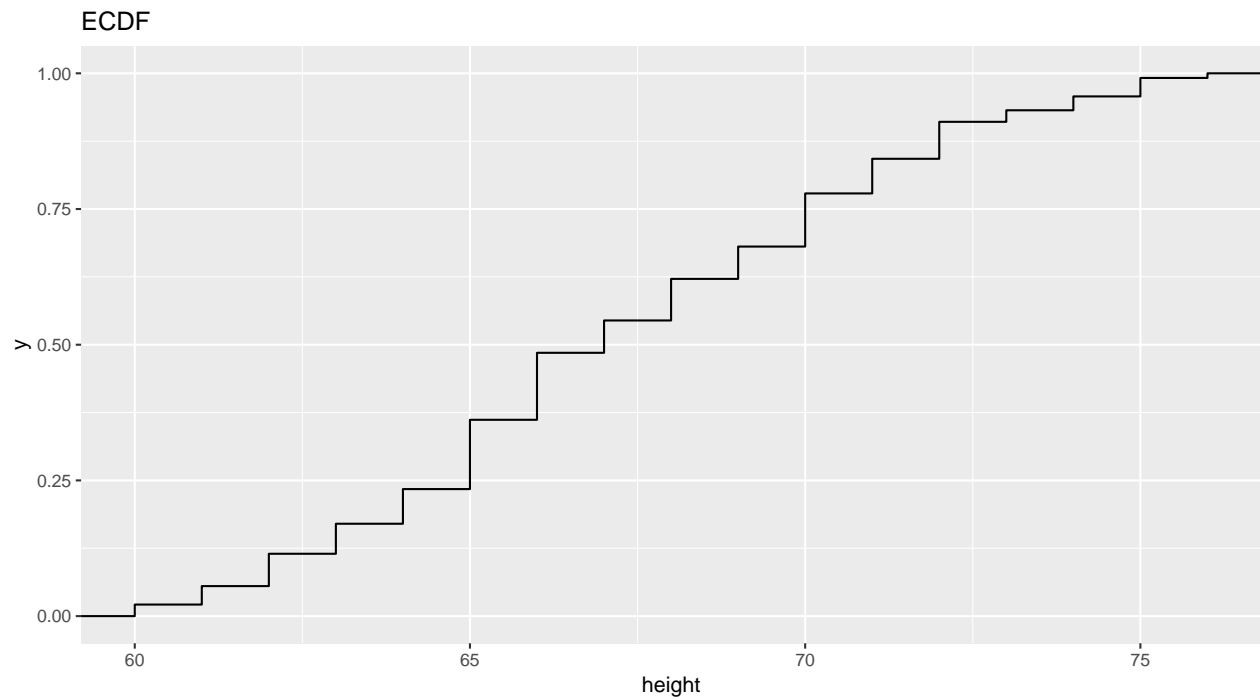


Basic Plots

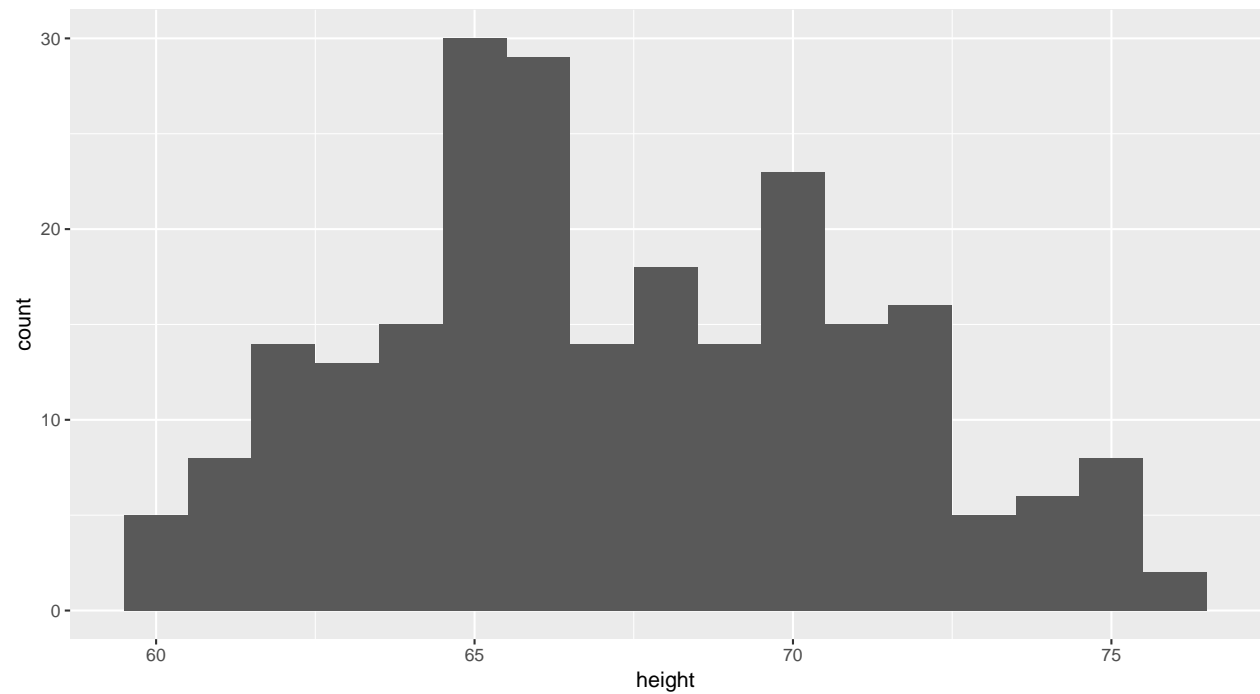
```
# Alternate way to get ECDF  
singer_gg <- ggplot(singer, aes(x = height))  
singer_gg + stat_ecdf()
```



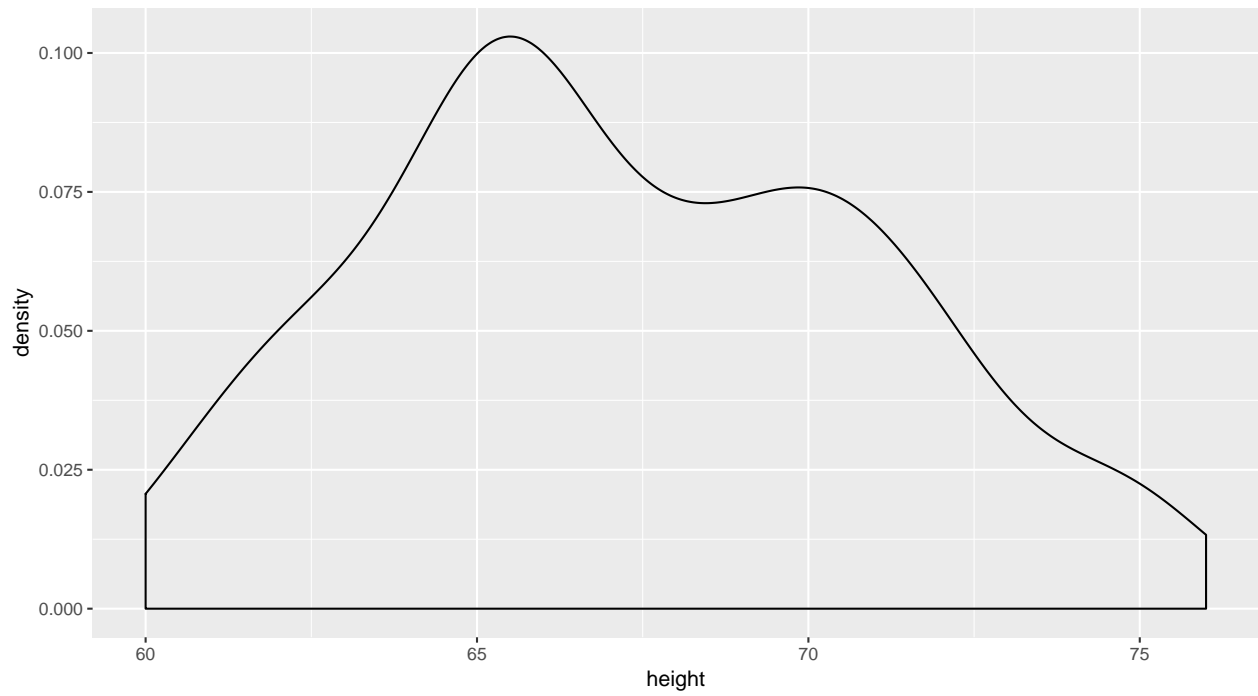
```
# Add Labels  
singer_gg + stat_ecdf() + ggtitle('ECDF')
```



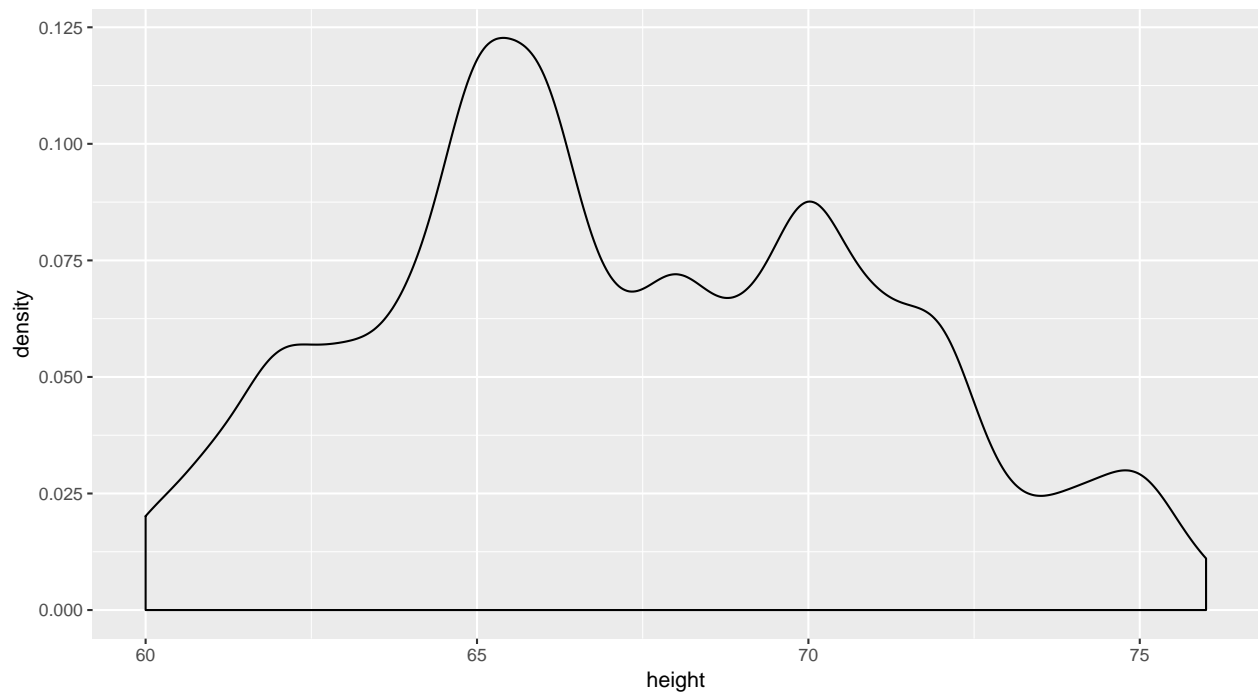
```
# Histogram  
#ggplot(singer, aes(x= height)) + geom_histogram()  
ggplot(singer, aes(x= height)) + geom_histogram(binwidth = 1)
```



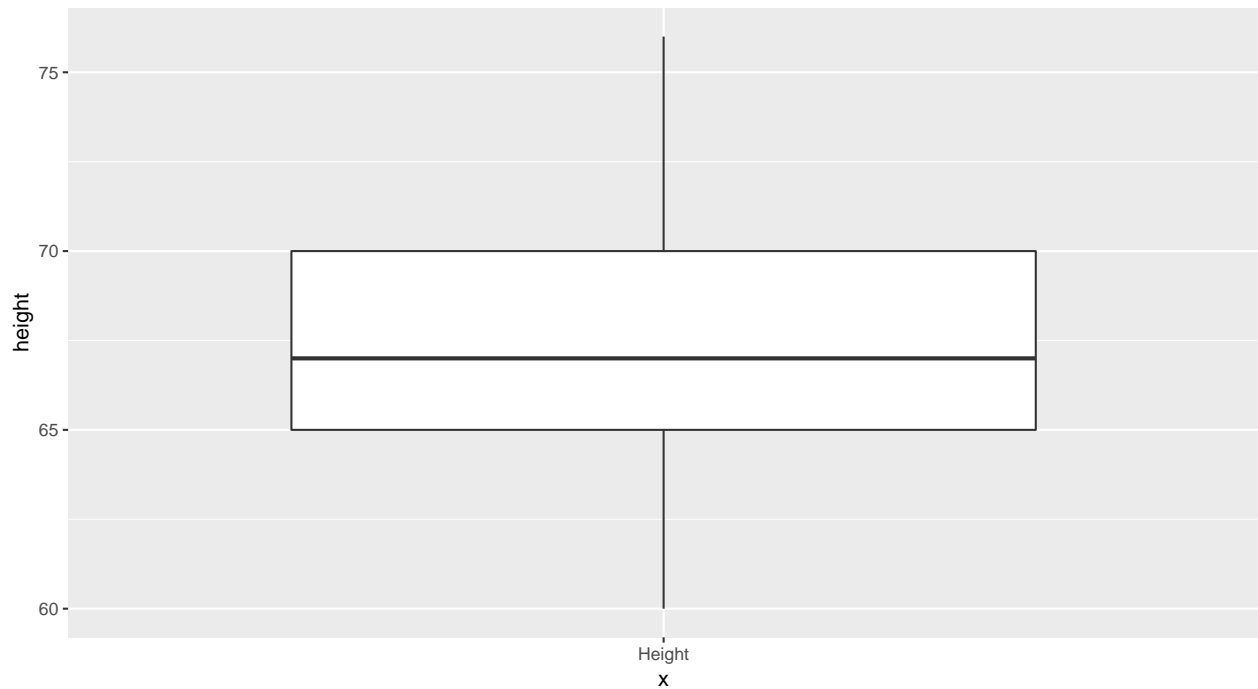
```
# Density Plot  
ggplot(singer, aes(x= height)) +  
  geom_density()
```



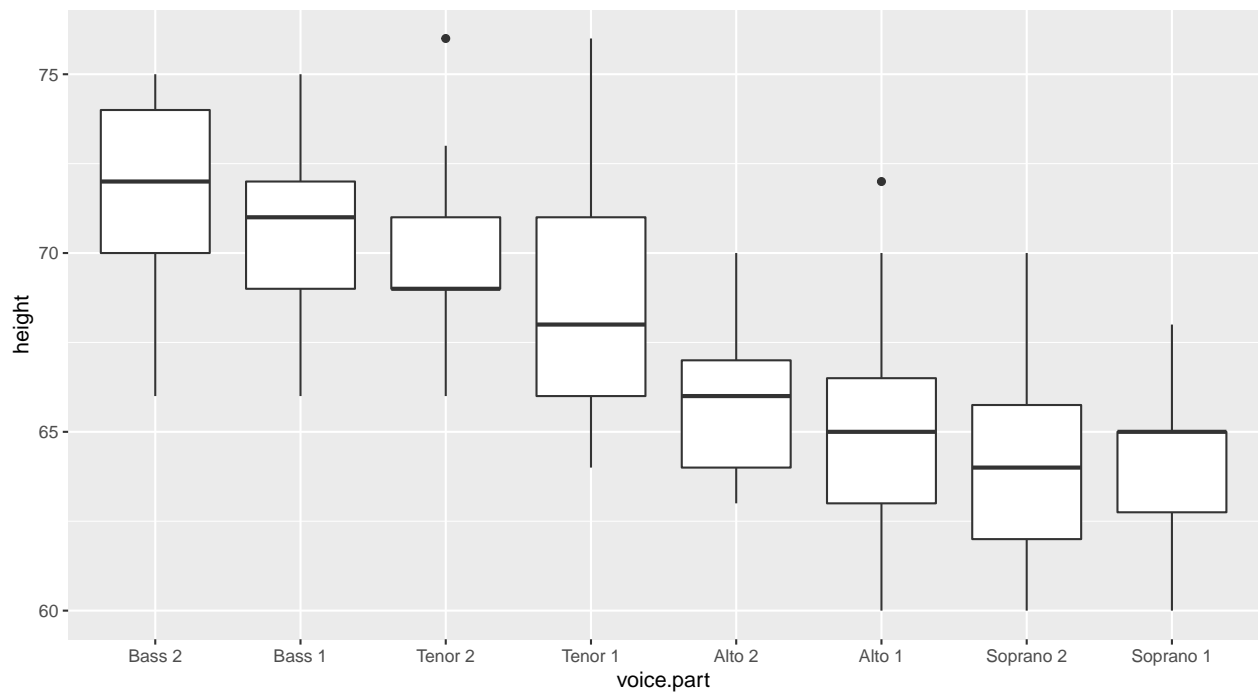
```
ggplot(singer, aes(x= height)) +  
  geom_density(adjust = 0.5)
```



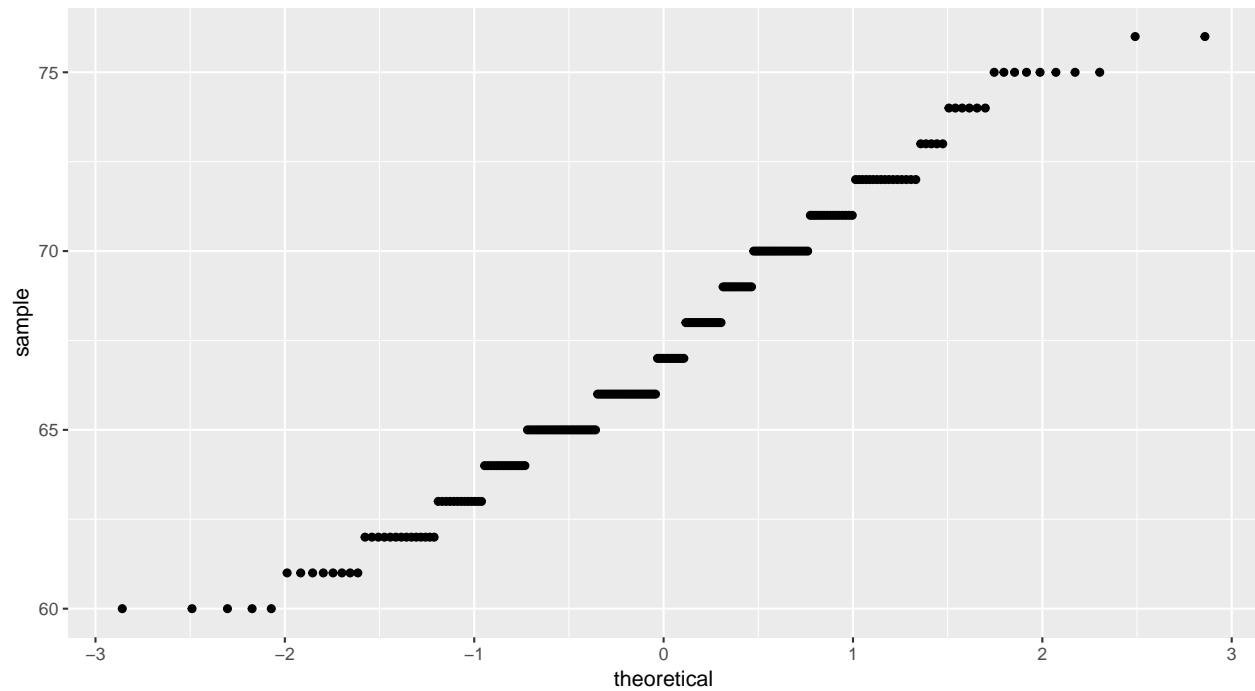
```
# Boxplot  
ggplot(singer, aes(x= "Height", y = height)) + geom_boxplot()
```



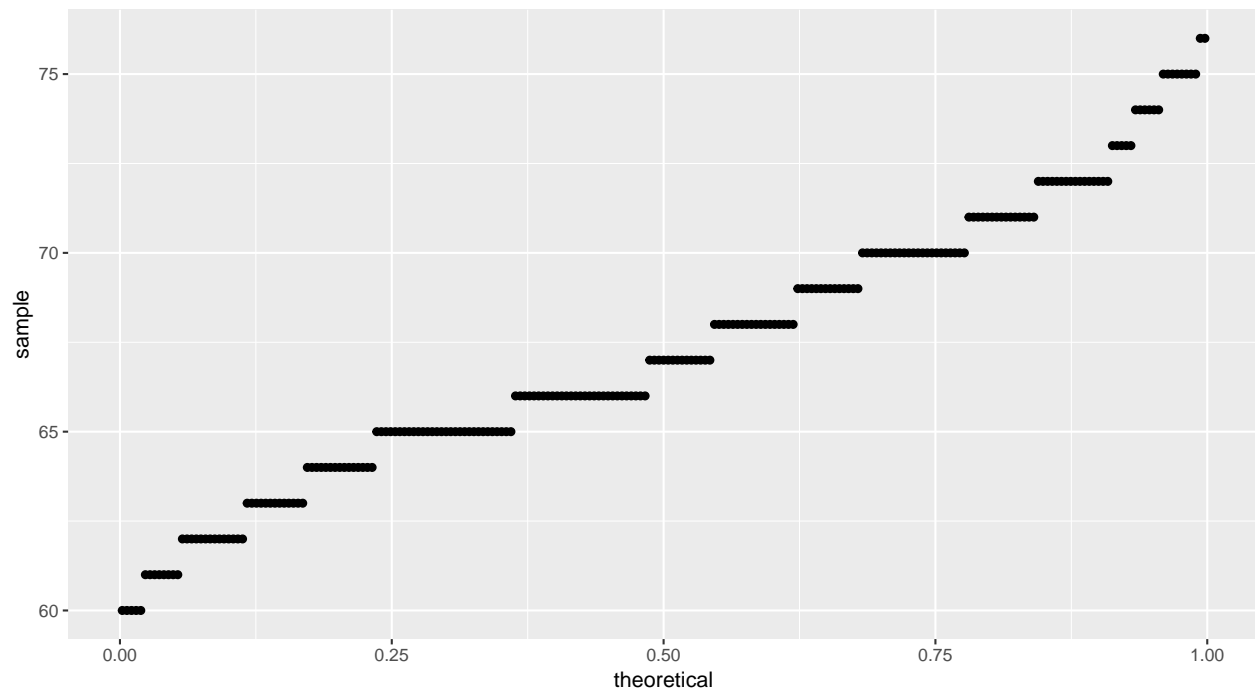
```
# Multiple Boxplots
ggplot(singer, aes(x = voice.part, y = height)) + geom_boxplot()
```



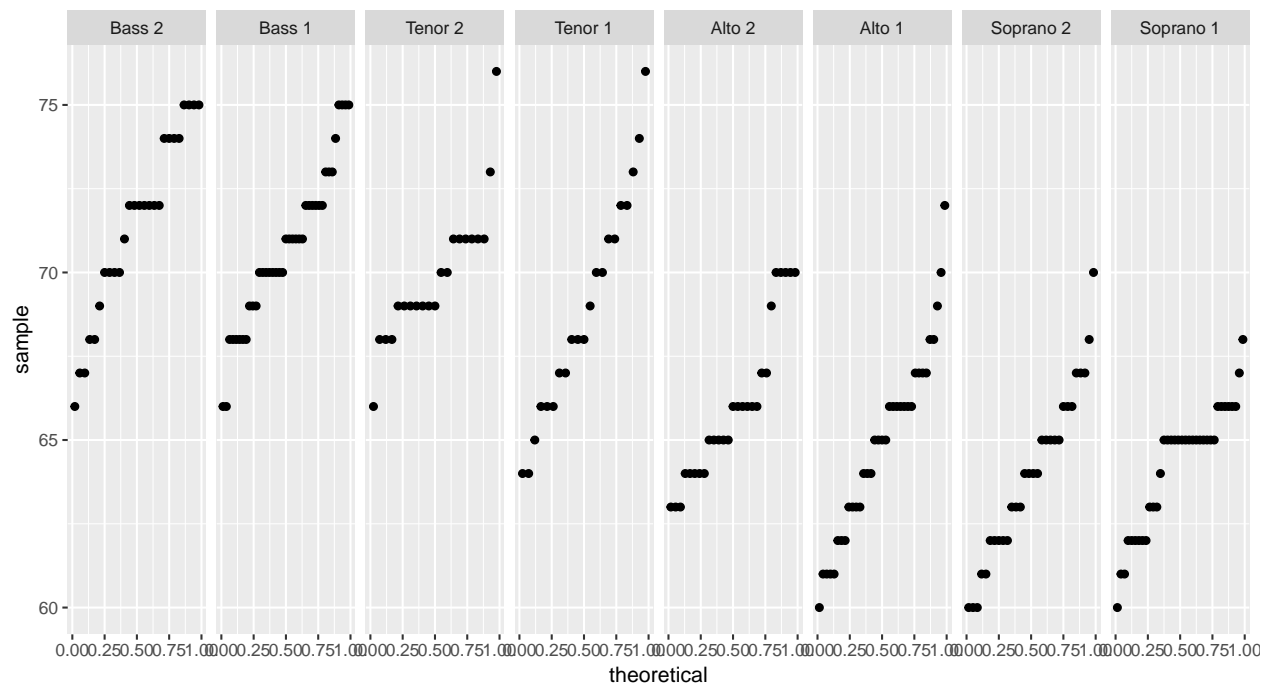
```
# QQ-Plot using a Normal Distribution
ggplot(singer, aes(sample = height)) + stat_qq()
```



```
# QQ-Plot using a Uniform Distribution
ggplot(singer, aes(sample = height)) +
  stat_qq(distribution = qunif)
```

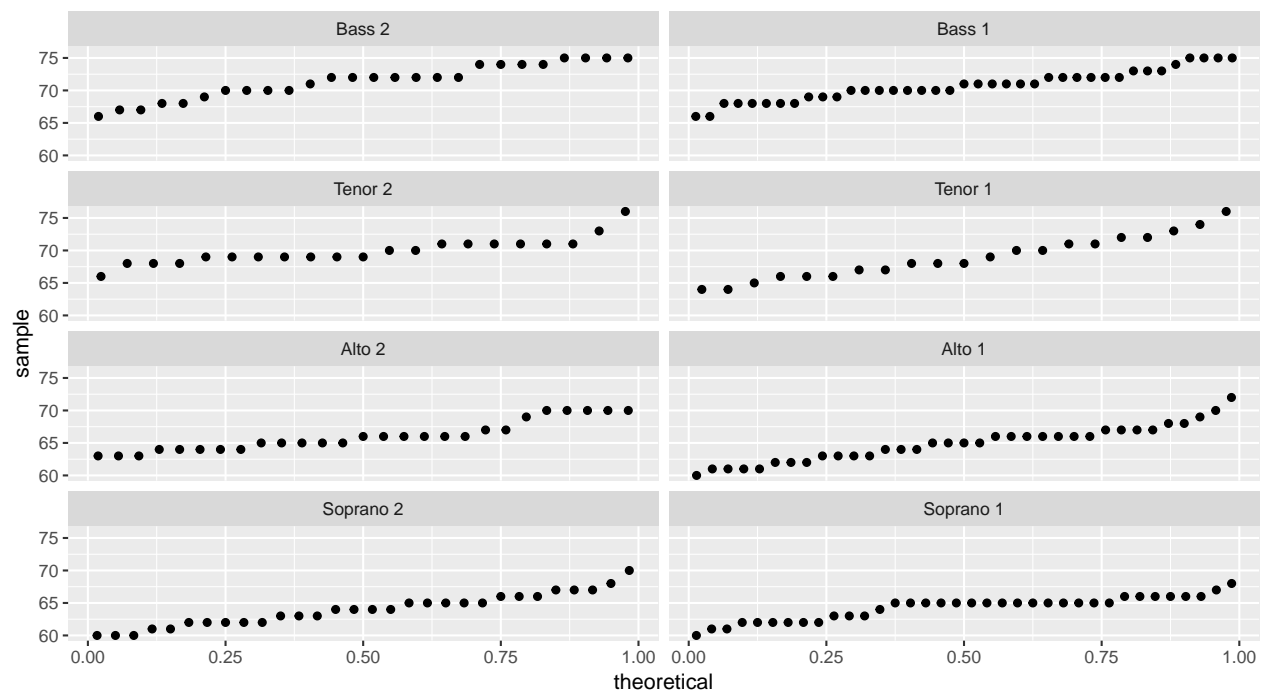


```
# Facet
ggplot(singer, aes(sample = height)) +
  stat_qq(distribution = qunif) +
  facet_grid(~voice.part)
```



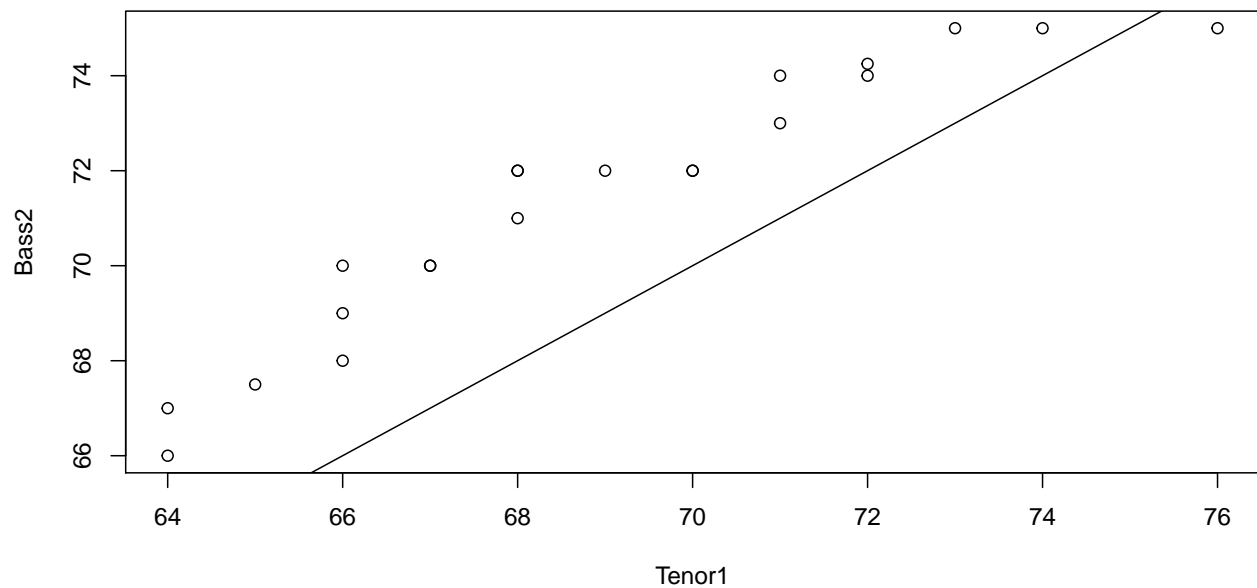
The plot above looks cramped, so let us try a grid with columns

```
ggplot(singer, aes(sample = height)) +  
  stat_qq(distribution = qunif) +  
  facet_wrap(~voice.part, ncol = 2)
```

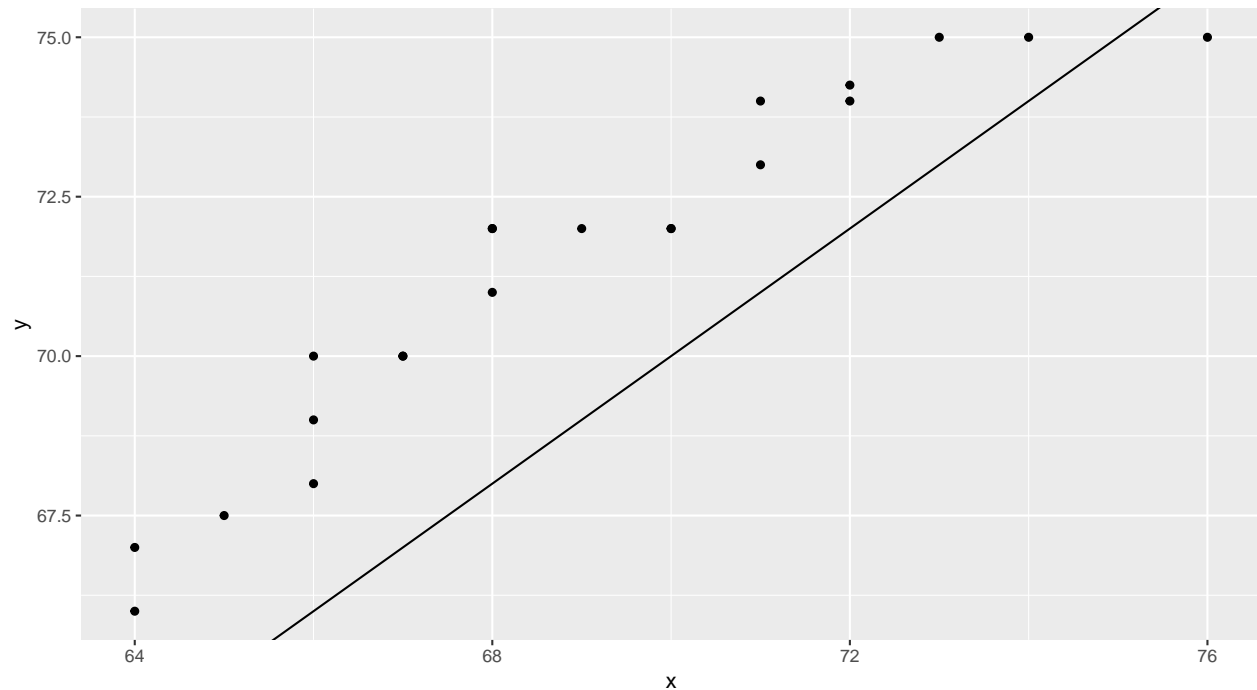


QQ-Plot using Base R function qqplot()

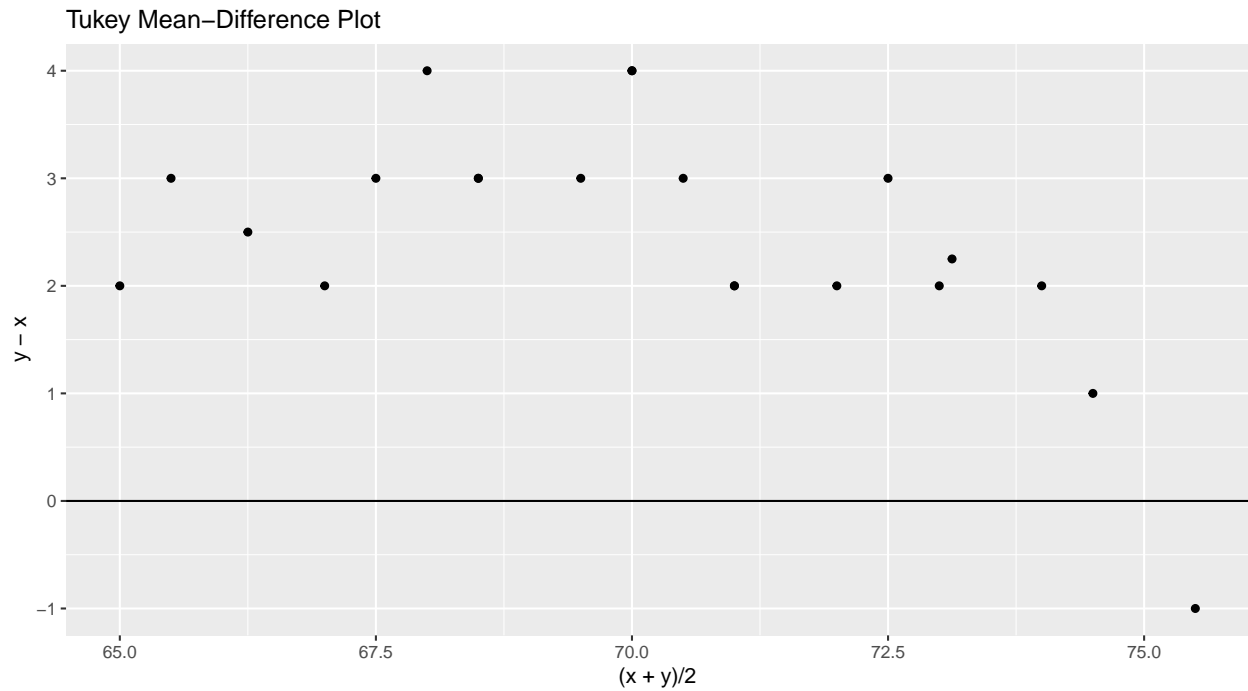
```
Tenor1 = singer$height[singer$voice.part == "Tenor 1"]  
Bass2 = singer$height[singer$voice.part == "Bass 2"]  
qqplot(Tenor1, Bass2)  
abline(0, 1)
```



```
# Using ggplot
#library(tibble)
#qq_df <- as.tibble(qqplot(Tenor1, Bass2, plot.it = FALSE))
qq_df <- as.data.frame(qqplot(Tenor1, Bass2, plot.it = FALSE))
ggplot(data = qq_df, mapping =
  aes(x = x, y = y)) +
  geom_point() +
  geom_abline()
```

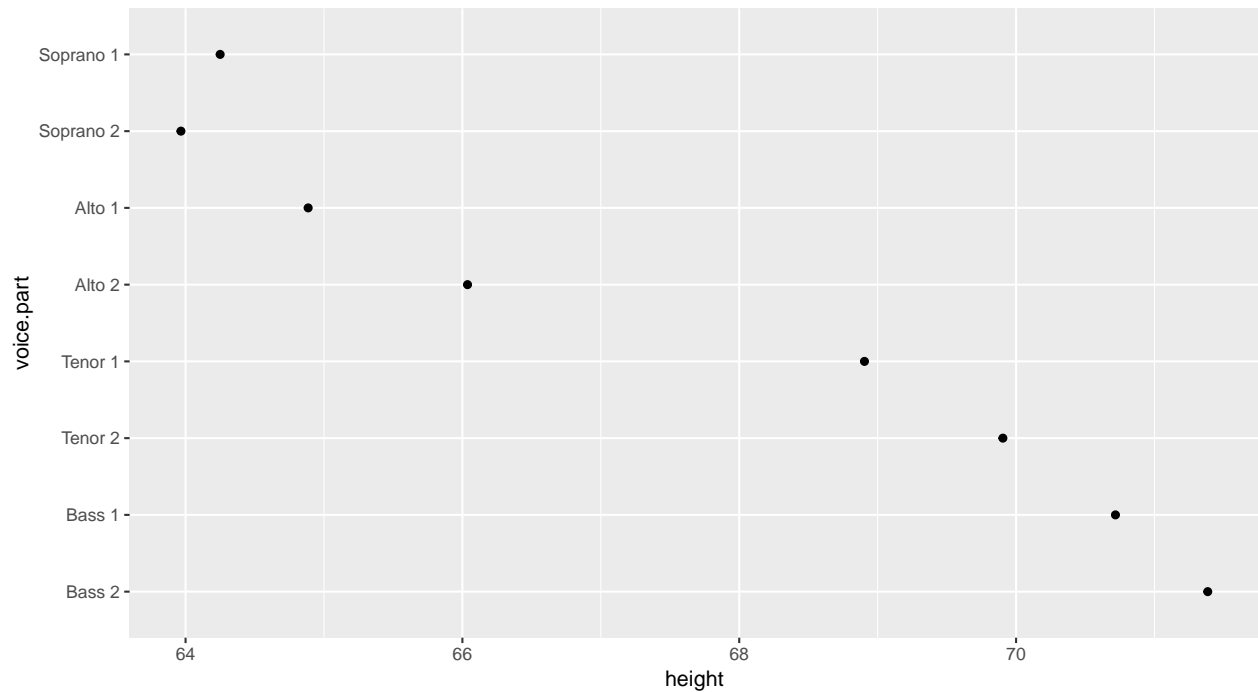


```
## Tukey-Mean difference Plot
ggplot(data = qq_df, mapping = aes(x = (x + y)/2, y = y - x)) +
  geom_point() +
  geom_abline(slope = 0) + ggtitle("Tukey Mean-Difference Plot")
```

```
singer_means <- aggregate(height ~ voice.part, FUN = mean, data = singer)
```

```
ggplot(singer_means, aes(x = voice.part, y = height)) + geom_point() + coord_flip()
```

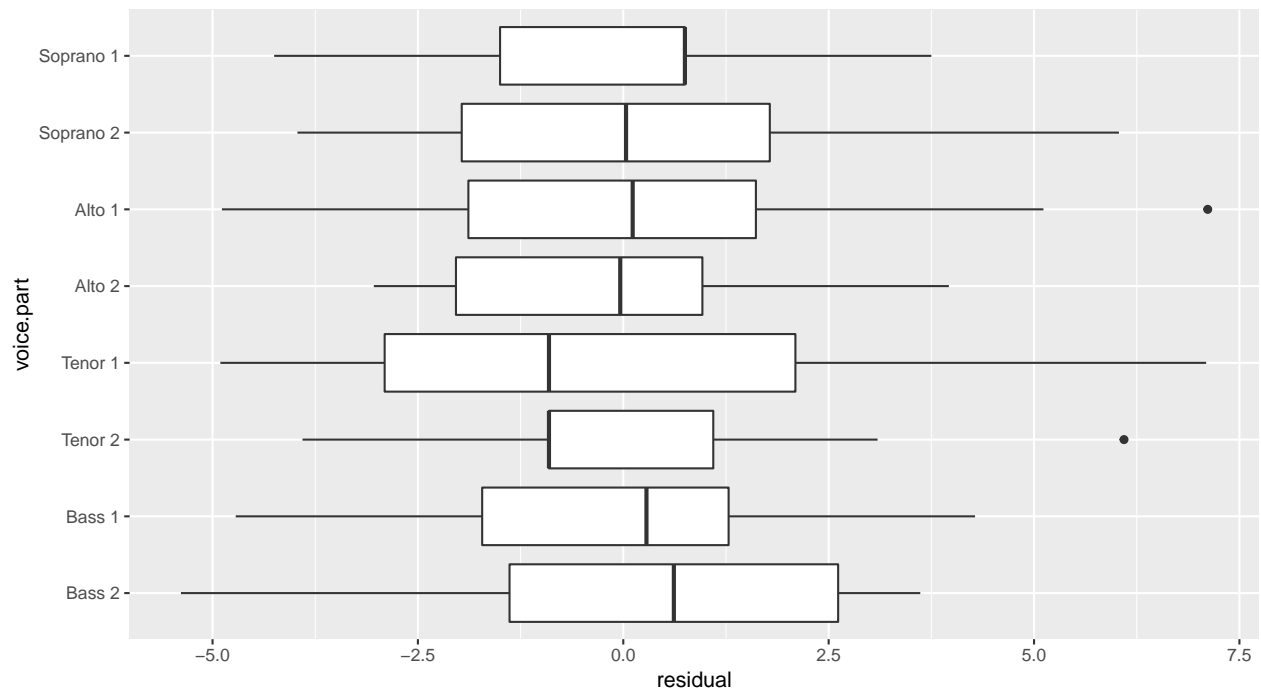


```
# Fitting a linear model
singer.lm <- lm(height ~ voice.part, data = singer)
```

```
# Extracting residual values
singer.res <- data.frame(voice.part = singer$voice.part, residual = residuals(singer.lm))
```

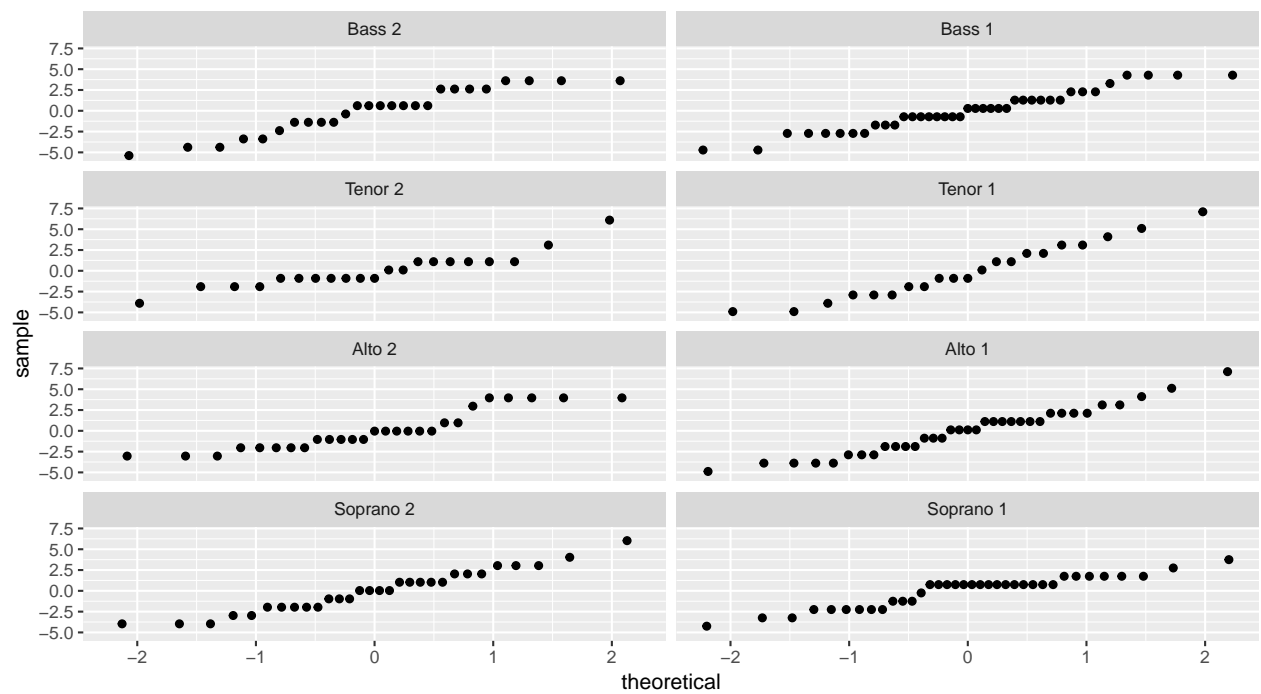
```
# Observing Residuals using boxplots
```

```
ggplot(singer.res, aes(x = voice.part, y = residual)) + geom_boxplot() + coord_flip()
```



```
# Checking normality of residuals
```

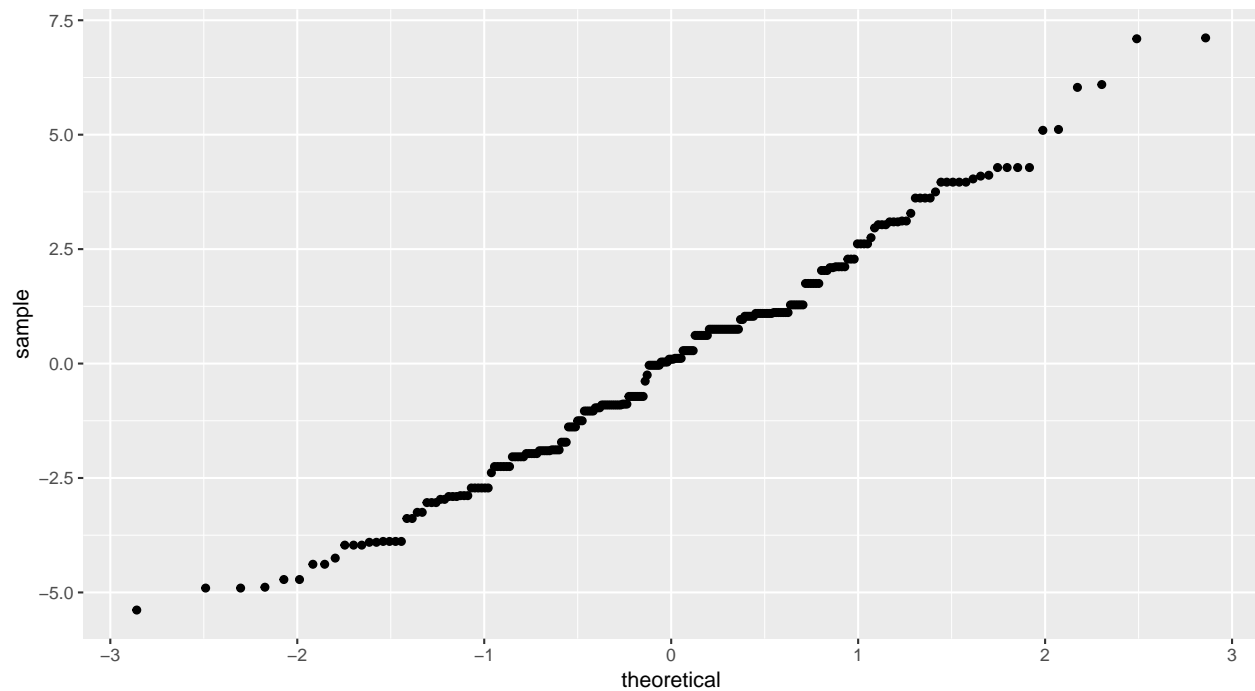
```
ggplot(singer.res, aes(sample = residual)) +  
  stat_qq() +  
  facet_wrap(~voice.part, ncol = 2)
```



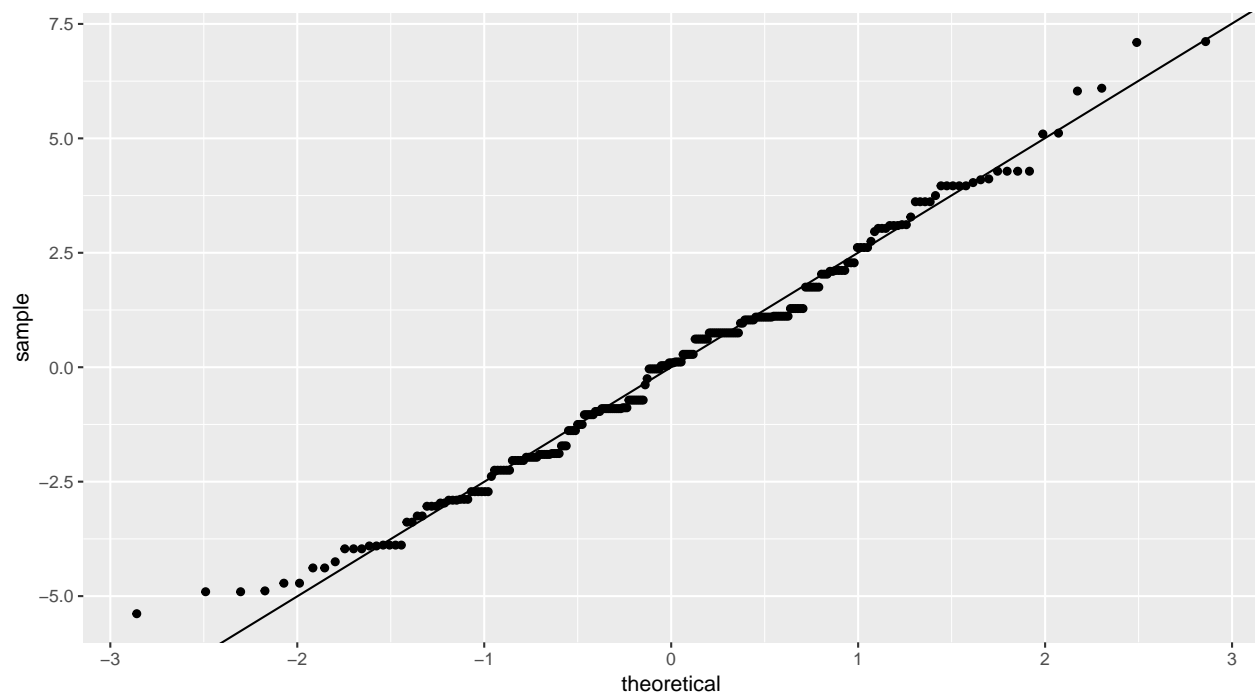
```
# QQ Plot for Normality
```

```
ggplot(singer.res, aes(sample = residual)) +
```

```
stat_qq()
```



```
# QQ Plot for Normality and adding slope and intercept  
ggplot(singer.res, aes(sample = residual)) +  
  stat_qq() + geom_abline(intercept = 0,  
slope = summary(singer.lm)$sigma)
```

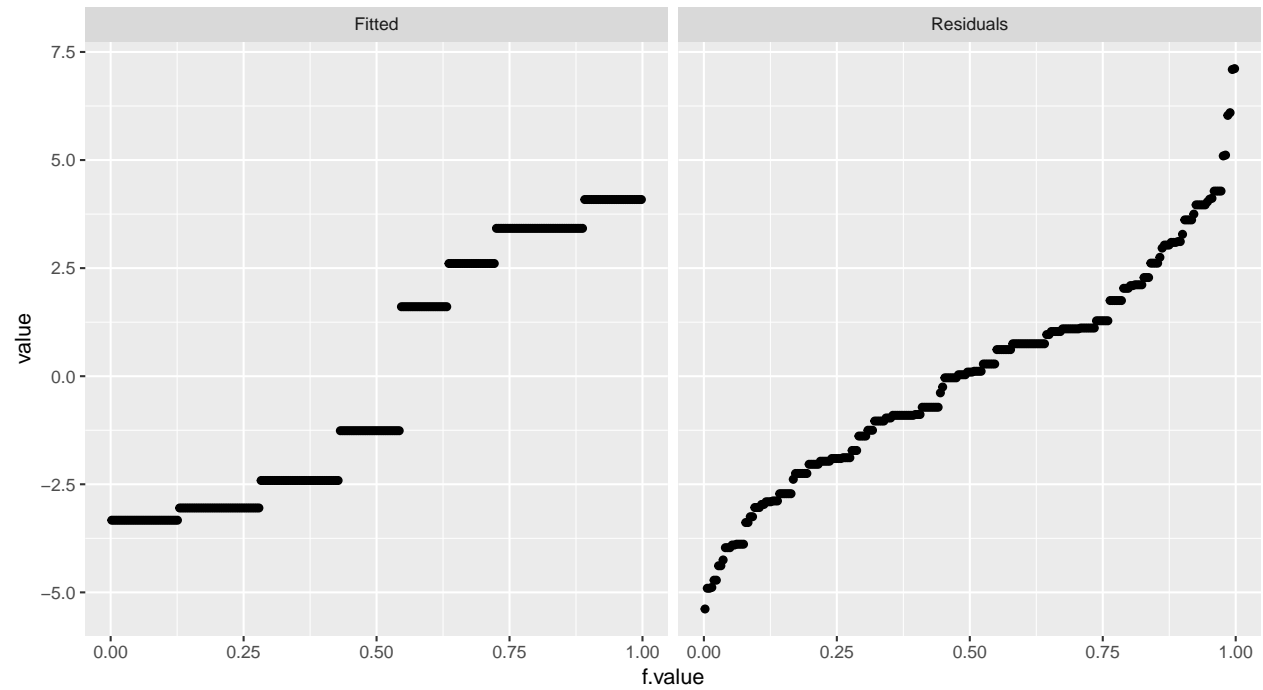


```
# Prepare Data  
singer.fitted <- sort(fitted.values(singer.lm)) - mean(fitted.values(singer.lm))  
singer.residuals <- sort(residuals(singer.lm))
```

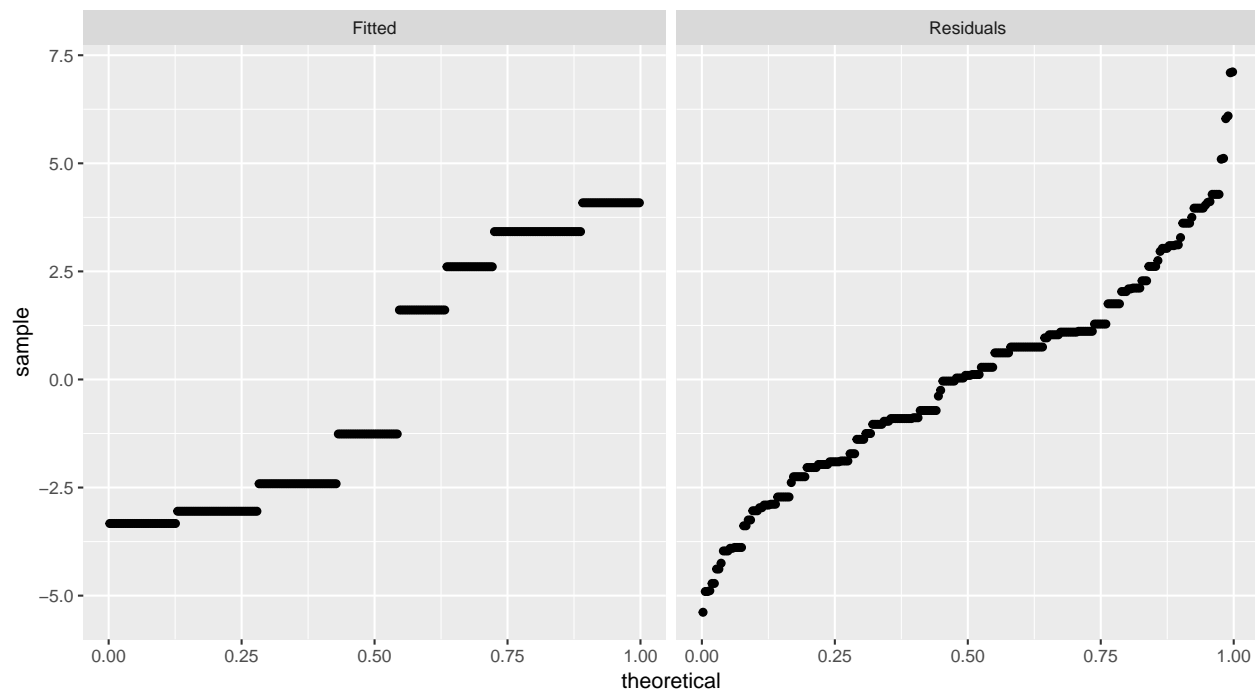
```
n = length(singer.residuals)
f.value = (0.5:(n - 0.5))/n
singer.fit = data.frame(f.value, Fitted = singer.fitted, Residuals = singer.residuals)
```

```
library(tidyr)
singer.fit.long = singer.fit %>% gather(type, value, Fitted:Residuals)
```

```
# Residual-Fit spread plot
ggplot(singer.fit.long, aes(x = f.value, y = value)) +
  geom_point() +
  facet_wrap(~type)
```



```
# Using stat_qq() to calculate f-values
ggplot(singer.fit.long, aes(sample = value)) +
  stat_qq(distribution = "qunif") +
  facet_grid(~type)
```



```
# Using R-Squared
```

```
round(var(singer.fitted)/var(singer$height),3)
```

```
## [1] 0.585
```

```
url <- "https://github.com/hadley/tidy-data/raw/master/data/billboard.csv"
```

```
billboard.raw <- read.csv(url, stringsAsFactors = FALSE)
```

```
# Creating a Tibble
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, setequal, union
```

```
billboard <- tbl_df(billboard.raw)
```

```
billboard
```

```
## # A tibble: 317 x 83
```

```
##   year artist.inverted track time genre date.entered date.peaked
##   <int> <chr>          <chr> <chr> <chr> <chr>      <chr>
## 1 2000 Destiny's Child Inde~ 3:38 Rock 2000-09-23 2000-11-18
## 2 2000 Santana        Mari~ 4:18 Rock 2000-02-12 2000-04-08
## 3 2000 Savage Garden I Kn~ 4:07 Rock 1999-10-23 2000-01-29
## 4 2000 Madonna        Music 3:45 Rock 2000-08-12 2000-09-16
## 5 2000 Aguilera, Chri~ Come~ 3:38 Rock 2000-08-05 2000-10-14
## 6 2000 Janet          Does~ 4:17 Rock 2000-06-17 2000-08-26
## 7 2000 Destiny's Child Say ~ 4:31 Rock 1999-12-25 2000-03-18
```

```
## 8 2000 Iglesias, Enri~ Be W~ 3:36 Latin 2000-04-01 2000-06-24
## 9 2000 Sisqo Inco~ 3:52 Rock 2000-06-24 2000-08-12
## 10 2000 Lonestar Amaz~ 4:25 Coun~ 1999-06-05 2000-03-04
## # ... with 307 more rows, and 76 more variables: x1st.week <int>,
## # x2nd.week <int>, x3rd.week <int>, x4th.week <int>, x5th.week <int>,
## # x6th.week <int>, x7th.week <int>, x8th.week <int>, x9th.week <int>,
## # x10th.week <int>, x11th.week <int>, x12th.week <int>,
## # x13th.week <int>, x14th.week <int>, x15th.week <int>,
## # x16th.week <int>, x17th.week <int>, x18th.week <int>,
## # x19th.week <int>, x20th.week <int>, x21st.week <int>,
## # x22nd.week <int>, x23rd.week <int>, x24th.week <int>,
## # x25th.week <int>, x26th.week <int>, x27th.week <int>,
## # x28th.week <int>, x29th.week <int>, x30th.week <int>,
## # x31st.week <int>, x32nd.week <int>, x33rd.week <int>,
## # x34th.week <int>, x35th.week <int>, x36th.week <int>,
## # x37th.week <int>, x38th.week <int>, x39th.week <int>,
## # x40th.week <int>, x41st.week <int>, x42nd.week <int>,
## # x43rd.week <int>, x44th.week <int>, x45th.week <int>,
## # x46th.week <int>, x47th.week <int>, x48th.week <int>,
## # x49th.week <int>, x50th.week <int>, x51st.week <int>,
## # x52nd.week <int>, x53rd.week <int>, x54th.week <int>,
## # x55th.week <int>, x56th.week <int>, x57th.week <int>,
## # x58th.week <int>, x59th.week <int>, x60th.week <int>,
## # x61st.week <int>, x62nd.week <int>, x63rd.week <int>,
## # x64th.week <int>, x65th.week <int>, x66th.week <lgl>,
## # x67th.week <lgl>, x68th.week <lgl>, x69th.week <lgl>,
## # x70th.week <lgl>, x71st.week <lgl>, x72nd.week <lgl>,
## # x73rd.week <lgl>, x74th.week <lgl>, x75th.week <lgl>, x76th.week <lgl>
```

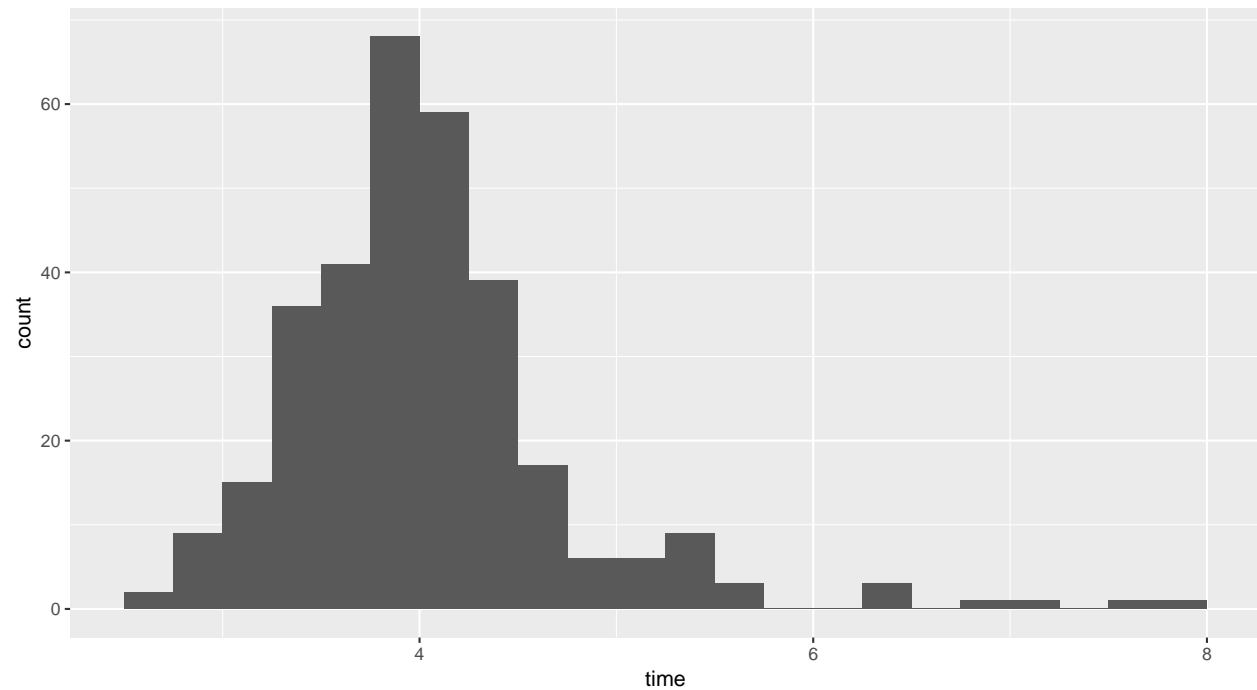
```
# Calculating time-length of each song
```

```
billboard.time <- strsplit(billboard$time, ":")
billboard.time <- matrix(unlist(billboard.time), byrow = T, ncol = 2)
billboard.mins <- as.numeric(billboard.time[, 1])
billboard.secs <- as.numeric(billboard.time[, 2])
billboard.time <- billboard.mins * 60 + billboard.secs
billboard$time <- billboard.time/60
```

```
billboard.long <- billboard %>%
  gather(week, rank, x1st.week:x76th.week, na.rm = TRUE)
billboard.long
```

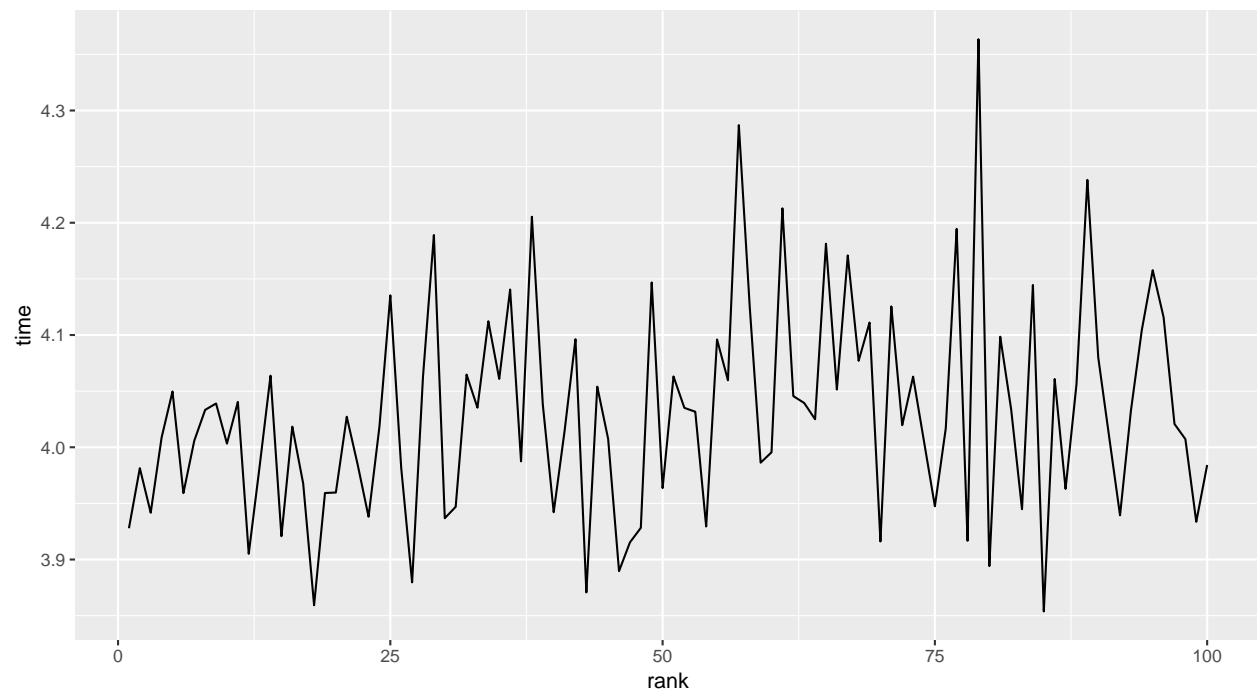
```
## # A tibble: 5,307 x 9
##   year artist.inverted track time genre date.entered date.peaked week
##   <int> <chr>          <chr> <dbl> <chr> <chr>          <chr>          <chr>
## 1 2000 Destiny's Child Inde~ 3.63 Rock 2000-09-23 2000-11-18 x1st~
## 2 2000 Santana Mari~ 4.3 Rock 2000-02-12 2000-04-08 x1st~
## 3 2000 Savage Garden I Kn~ 4.12 Rock 1999-10-23 2000-01-29 x1st~
## 4 2000 Madonna Music 3.75 Rock 2000-08-12 2000-09-16 x1st~
## 5 2000 Aguilera, Chri~ Come~ 3.63 Rock 2000-08-05 2000-10-14 x1st~
## 6 2000 Janet Does~ 4.28 Rock 2000-06-17 2000-08-26 x1st~
## 7 2000 Destiny's Child Say ~ 4.52 Rock 1999-12-25 2000-03-18 x1st~
## 8 2000 Iglesias, Enri~ Be W~ 3.6 Latin 2000-04-01 2000-06-24 x1st~
## 9 2000 Sisqo Inco~ 3.87 Rock 2000-06-24 2000-08-12 x1st~
## 10 2000 Lonestar Amaz~ 4.42 Coun~ 1999-06-05 2000-03-04 x1st~
## # ... with 5,297 more rows, and 1 more variable: rank <int>
```

```
ggplot(billboard, aes(x = time)) +  
  geom_histogram(breaks = seq(2.5, 8, 0.25))
```



```
# Finding mean song length  
time.means = aggregate(time ~ rank, FUN = mean, data = billboard.long)
```

```
# Plotting mean song lengths  
ggplot(time.means, aes(x = rank, y = time)) + geom_line()
```



Transformations

Log Transformation

```
# Load data from workspace  
load("../data/lattice.RData")
```


```
summary(fusion.time)
```

```
##      time      nv.vv  
## Min.   : 1.000    NV:43  
## 1st Qu.: 2.475    VV:35  
## Median : 5.250  
## Mean   : 7.210  
## 3rd Qu.: 9.050  
## Max.   :47.200
```

```
# Plot without data transformation
```

```
time = fusion.time$time  
nv.vv = fusion.time$nv.vv  
NV.times = sort(time[nv.vv == "NV"])  
VV.times = sort(time[nv.vv == "VV"])  
NV.VV.qq = as.data.frame(qqplot(NV.times, VV.times, plot.it = FALSE))
```


```
ggplot(NV.VV.qq, aes(x, y)) + geom_point() + geom_abline()
```




```
# Using a log transformation
```

```
NV.times.log = sort(log2(time[nv.vv == "NV"]))  
VV.times.log = sort(log2(time[nv.vv == "VV"]))  
NV.VV.qq.log = as.data.frame(qqplot(NV.times.log, VV.times.log, plot.it = FALSE))
```

```
ggplot(NV.VV.qq.log, aes(x, y)) + geom_point() + geom_abline()
```

```
# Tukey Mean-Difference Plot
ggplot(NV.VV.qq.log, aes((x + y)/2, y - x)) + geom_point()
```




Power Transformation

```
n.VV = length(VV.times)
power = rep(seq(-1, 1, 0.25), each = n.VV)
VV.time = c(VV.times^-1, VV.times^-0.75,
            VV.times^-0.5, VV.times^-0.25,
```

```
log(VV.times), VV.times^0.25,
VV.times^0.5, VV.times^0.75, VV.times)
```


```
ggplot(data.frame(power, VV.time), aes(sample = VV.time)) +
  stat_qq() + facet_wrap(~power, scales = "free")
```




```
summary(food.web)
```

```
## mean.length dimension
## Min. :1.670 Two :40
## 1st Qu.:2.220 Mixed:45
## Median :2.710 Three:28
## Mean :2.885
## 3rd Qu.:3.350
## Max. :5.920
```

```
ggplot(food.web, aes(sample = mean.length)) +
  stat_qq(distribution = qunif) + facet_grid(~dimension)
```




```
web.length = food.web$mean.length
dimension = food.web$dimension
n = nrow(food.web)
median.3 = median(web.length[dimension == "Three"])
median.2 = median(web.length[dimension == "Two"])
median.mixed = median(web.length[dimension == "Mixed"])
group.median = rep(NA, n)
group.median[dimension == "Three"] = median.3
group.median[dimension == "Two"] = median.2
group.median[dimension == "Mixed"] = median.mixed
jittered.medians = group.median + runif(n, -0.1, 0.1)
root.abs.res = sqrt(abs(web.length - group.median))
food.web.sl = data.frame(jittered.medians, root.abs.res, dimension)
ggplot(food.web.sl, aes(jittered.medians, root.abs.res, col = dimension)) +
  geom_point()
```




Using log transformation

```
log.web.length = log2(food.web$mean.length)
median.3.log = median(log.web.length[dimension == "Three"])
median.2.log = median(log.web.length[dimension == "Two"])
median.mixed.log = median(log.web.length[dimension == "Mixed"])
group.median.log = rep(NA, n)
group.median.log[dimension == "Three"] = median.3.log
group.median.log[dimension == "Two"] = median.2.log
group.median.log[dimension == "Mixed"] = median.mixed.log
jittered.medians.log = group.median.log + runif(n, -0.05, 0.05)
root.abs.res.log = sqrt(abs(log.web.length - group.median.log))
food.web.log.sl = data.frame(jittered.medians.log, root.abs.res.log, dimension)
ggplot(food.web.log.sl, aes(jittered.medians.log, root.abs.res.log, col = dimension)) +
  geom_point()
```



```
ggplot(food.web, aes(sample = log(mean.length))) +
  stat_qq() +
  facet_wrap(~dimension)
```




```
# Using Inverse Transformation
inv.web.length = 1/food.web$mean.length
median.3.inv = median(inv.web.length[dimension == "Three"])
median.2.inv = median(inv.web.length[dimension == "Two"])
median.mixed.inv = median(inv.web.length[dimension == "Mixed"])
group.median.inv = rep(NA, n)
```

```

group.median.inv[dimension == "Three"] = median.3.inv
group.median.inv[dimension == "Two"] = median.2.inv
group.median.inv[dimension == "Mixed"] = median.mixed.inv
jittered.medians.inv = group.median.inv + runif(n, -0.01, 0.01)
root.abs.res.inv = sqrt(abs(inv.web.length - group.median.inv))
food.web.inv.sl = data.frame(jittered.medians.inv, root.abs.res.inv, dimension)
ggplot(food.web.inv.sl, aes(jittered.medians.inv, root.abs.res.inv, col = dimension)) +
  geom_point()

```



```

# Calculate Mean
aggregate(root.abs.res.inv ~ dimension, FUN = mean)

```

```


##   dimension root.abs.res.inv
## 1      Two      0.2332382
## 2    Mixed      0.2563367
## 3     Three      0.2579944

```


```

# Normality Check
ggplot(food.web, aes(sample = 1/mean.length)) + stat_qq() + facet_wrap(~dimension)

```



```
food.web.lm = lm(inv.web.length ~ dimension)
food.web.res = residuals(food.web.lm)
res.qq.3 = qqplot(food.web.res, food.web.res[dimension == "Three"], plot.it = FALSE)
res.qq.2 = qqplot(food.web.res, food.web.res[dimension == "Two"], plot.it = FALSE)
res.qq.mixed = qqplot(food.web.res, food.web.res[dimension == "Mixed"], plot.it = FALSE)
food.web.res.qq = data.frame(pooled = c(res.qq.3$x, res.qq.2$x, res.qq.mixed$x),
  residual = c(res.qq.3$y, res.qq.2$y, res.qq.mixed$y), dimension = c(rep("Three",
    length(res.qq.3$x)), rep("Two", length(res.qq.2$x)), rep("Mixed", length(res.qq.mixed$x))))
ggplot(food.web.res.qq, aes(pooled, residual)) + geom_point() + geom_abline() +
  facet_wrap(~dimension)
```



```

# Residual Fitted Spread Plot
food.web.fitted = sort(fitted.values(food.web.lm)) - mean(fitted.values(food.web.lm))
n = length(inv.web.length)
f.value = (0.5:(n - 0.5))/n
food.web.fit = data.frame(f.value, Fitted = food.web.fitted, Residuals = sort(food.web.res))
library(tidyr)
food.web.fit.long = food.web.fit %>% gather(type, value, Fitted:Residuals)
ggplot(food.web.fit.long, aes(x = f.value, y = value)) + geom_point() + facet_wrap(~type)

```

